

ANALYSIS OF BIG DATA ARCHITECTURE

Dimensions of Big Data architectures

- **Data model(s):**
 - Relations, trees (XML, JSON), graphs (RDF, others...), nested relations
 - Query language
- **Heterogeneity** (DM, QL): none, some, a lot
- **Scale:** small (~10-20 sites) or large (~10.000 sites)
- **ACID** properties
- **Control:**
 - Single master w/complete control over N slaves (Hadoop/HDFS)
 - Sites publish independently and process queries as directed by single master/*mediator*
 - Many-mediator systems, or peer-to-peer (P2P) with *super-peers*
 - Sites completely independent (P2P)

对等网络

Architectures we will cover

- Distributed databases
- Mediator (data integration) systems
- Peer-to-peer data management systems
- Structured data management on top of MapReduce
- Dataspaces, polystores, datalakes


DISTRIBUTED RELATIONAL DATABASES

Distributed relational databases


- Oldest distributed architecture ('70s): IBM System R*
- Illustrate/introduce the main principles
- **Data** is distributed among many *nodes* (*sites, peers...*)
 - **Data catalog**: information on which data is stored where
 - Explicit : « All Paris sales are stored in Paris ».
Horizontal/vertical table fragmentation
Catalog stored at a master/central server.
 - Implicit: « Data is distributed by the value of the city »
(« somewhere »)
- **Queries** are distributed (may come from any site)
- **Query processing** is distributed
 - Operators may run on different sites → network transfer
 - Another layer of complexity to the optimization process

Distributed query optimization

Example 1: $R@s_1$, $S@s_2$, $T@s_3$, $q@s_4$



Example 2: $R@s_1$, $S@s_2$, $T@s_3$, $U@s_4$, $q@s_5$




Plan pruning criteria if all the sites and network connections have equal performance:


- Ship the smaller collection

Distributed query optimization

Example 1: $R@s1$, $S@s2$, $T@s3$, $q@s4$



Example 2: $R@s1$, $S@s2$, $T@s3$, $U@s4$, $q@s5$




Plan pruning criteria if all the sites and network connections have equal performance:


- Ship the smaller collection
- Transfer to join partner or the query site

Distributed query optimization

Example 1: $R@s1$, $S@s2$, $T@s3$, $q@s4$



Example 2: $R@s1$, $S@s2$, $T@s3$, $U@s4$, $q@s5$




Plan pruning criteria if all the sites and network connections have equal performance:

- Ship the smaller collection.
- Transfer to join partner or the query site

This plan illustrates total effort != response time

Distributed query optimization technique: semijoin reducers

- $R \text{ join } S = (R \text{ semijoin } S) \text{ join } S$



- Useful in distributed settings to reduce transfers: *if the distinct $S.b$ values are smaller than the non-matching R tuples*
- Symmetrical alternative: $R \text{ join } S = R \text{ join } (S \text{ semijoin } R)$
- This gives one more alternative in every join \rightarrow search space explosion
- Heuristics [Stocker, Kossmann et al., ICDE 2001]

Distribution of control in distributed relational DBs ('70s)

Servers DB1@site1: R1(a,b), S1(a,c)

Server DB2@site2: R2(a,b), S2(a,c),

Server DB3@site3: R3(a,b),
S3(a,c) defined as:

```
select * from DB1.S1 union all
```

```
select * from DB2.S2 union all
```

```
select R1.a as a, R2.b as c from DB1.R1 r1, DB2.R2 r2  
where r1.a=r2.a
```

DB3@site3 decides what to import from site1, site2 (« hard links »)

Site1, site2 are independent servers

Also: replication policies, distribution etc. (usually with one or a few masters)

Modern distributed databases: H-Store (subsequently VoltDB)

- From the team of Michael Stonebraker (Turing Award, author of the Postgres system)
- Main goal: quick OLTP (**o**nline **t**ransaction **p**rocessing), e.g., sales, likes, posts...
- Built to run on **cluster** for horizontal scalability
- **Share-nothing architecture**: each node stores tables **shards** (+ k replication for durability)

Frequent concept in Big Data architectures: shards


- **Shard** = small fragment of a data collection (e.g., a table)
- The assignment of data items (e.g. tuples) into shards is often done by **hashing** on tuple key
 - The table must have at least one key
 - Hashing ensures (with high probability) uniform distribution
- Key-based hashing is used as a mechanism for implementing distributed data catalogs. We will encounter it often.



Modern distributed RDB: MemSQL

MemSQL runs with

- a **master aggregator**, responsible of the metadata (catalog) 聚合器
- possibly more aggregators
- at least one **leaf**, each of which stores part(s) of some table(s)
- In each leaf, there are **partitions** (by default: 1 per CPU core)



Availability group: a set of machines + a set of replica machines (one-to-one)

Query processing in MemSQL

- **Indexes** managed within each partition
- In general, every query is run with a level of parallelism equal to the number of partitions
- **Select** queries are executed by the leaves which hold some partition(s) with data matching the query
- **Aggregation** queries run at the leaves involved and at the aggregator(s)
- **Join** queries
 - Easy if one input is a *reference* (small) table: one that is replicated fully to every machine in the cluster
 - Otherwise, they recommend sharing the shard key across tables to be joined
 - Otherwise, joins will incur data transfer within the cluster.

H-Store transactions

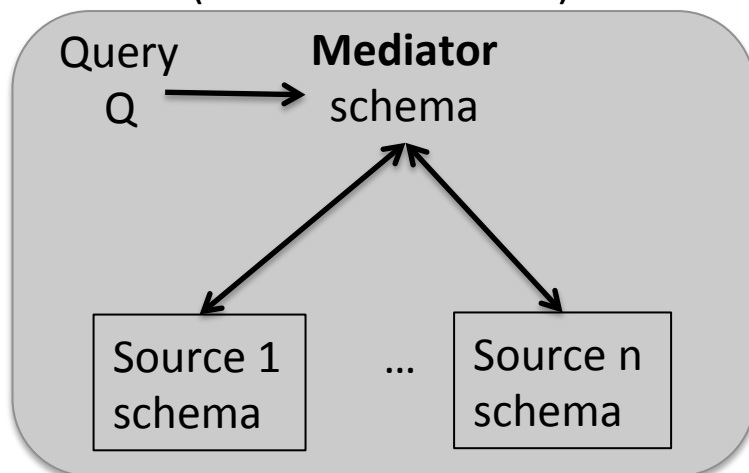
- Applications call **stored procedures** = code which also contains SQL queries
 - Each contained SQL query is partially unknown (depends on parameters specified at runtime); H-Store "pre-optimizes" it
- 1 **transaction** = 1 call of a stored procedure
- Can be submitted to any node (together with parameters)
- The node can run the procedure up to the query(ies) → updated, completely known plan → transaction manager

MEDIATOR SYSTEMS


Mediator systems

- A set of **data sources**, of the same or different data model, query language; source schemas
- A **mediator** data model and mediator schema
- Queries are asked against the mediator schema

Common data model
(sources+mediator)




Mediator data model



- **ACID:** mostly read-only; **size:** small
- **Control:** Independent publishing; mediator-driven integration

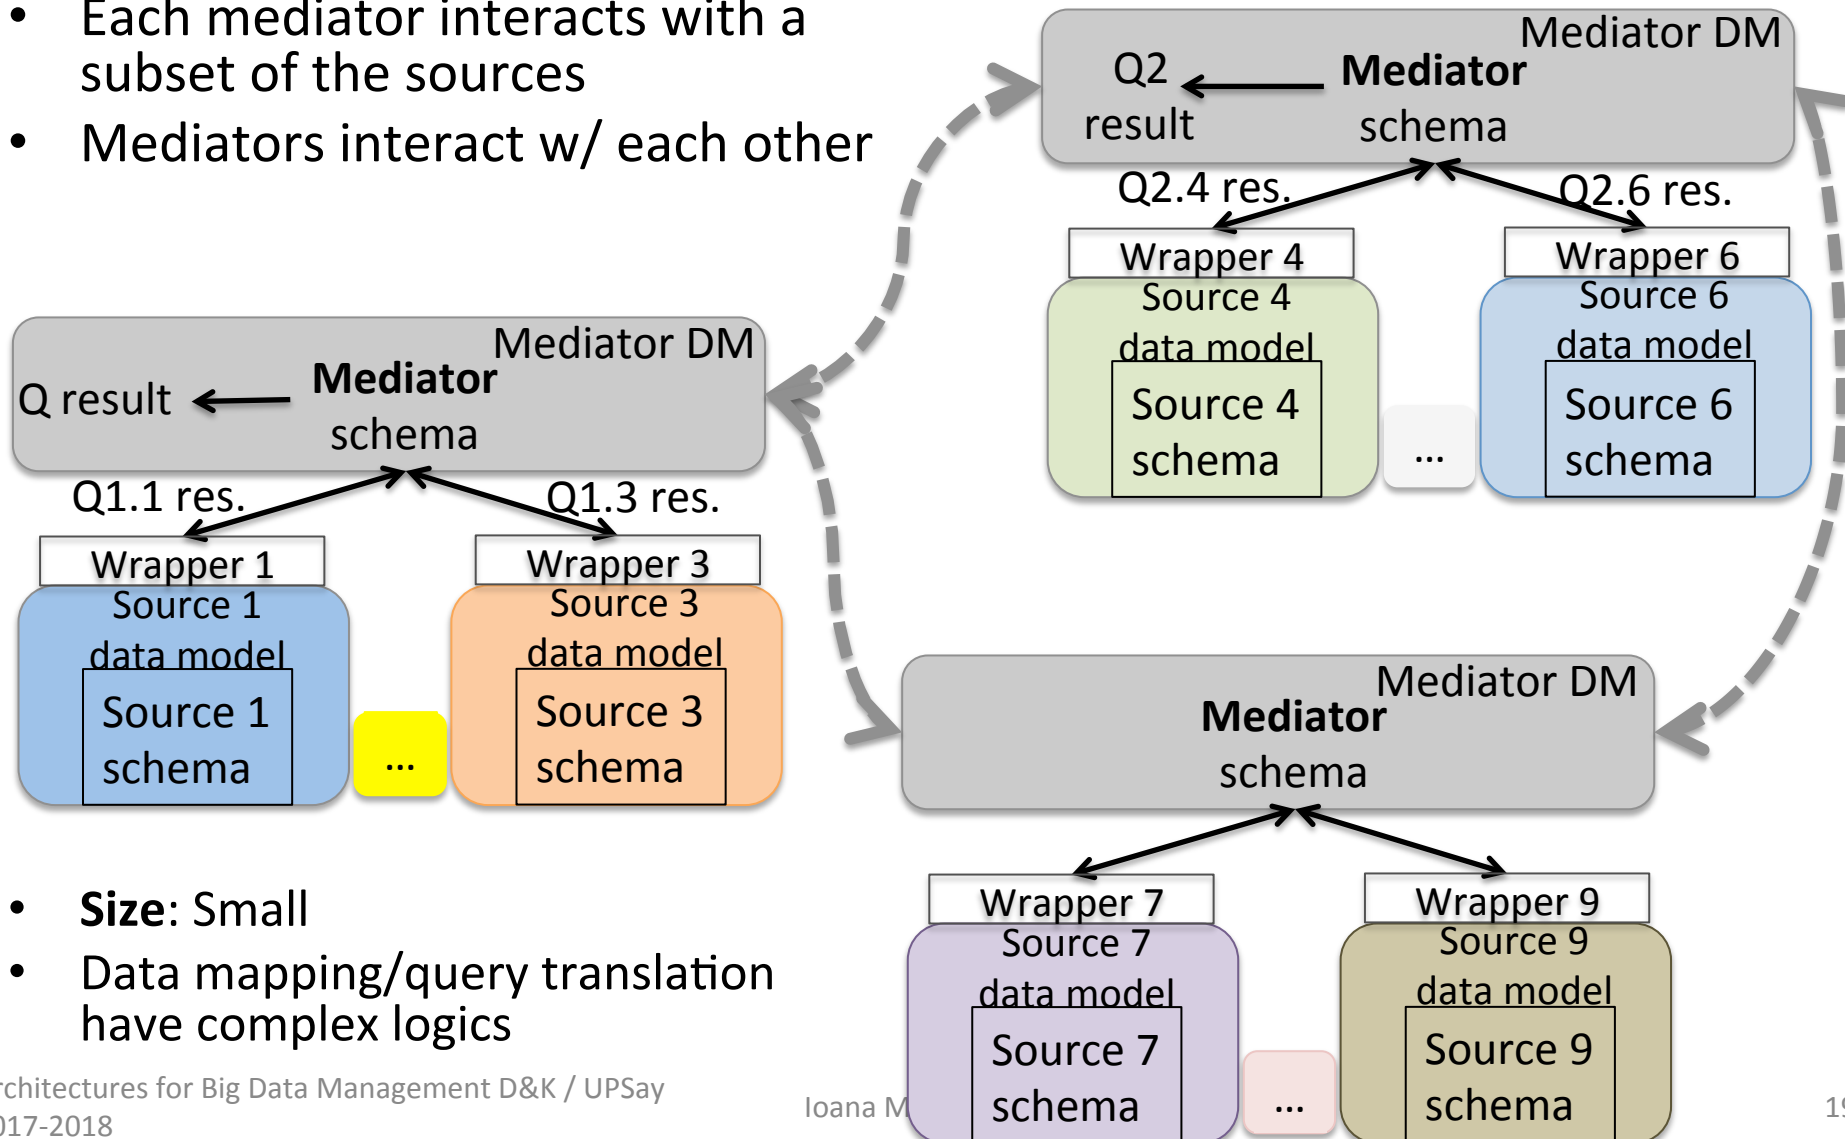
Many-mediator systems

- Each mediator interacts with a subset of the sources
- Mediators interact w/ each other



Many-mediator systems

- Each mediator interacts with a subset of the sources
- Mediators interact w/ each other



- **Size:** Small
- Data mapping/query translation have complex logics

Connecting the source schemas to the global schema

Example scenario:

Source **s1** has the schema:

ParisHotels(street, name, roomPrice)

Source **s2** has the schema:

LyonHotel(street, name, roomDesc, roomPrice)

Source **s3** has the schema:

Restaurants(city, street, name, rating)

The **global schema** is:

Hotel(city, street, name, descr, price)

Restaurants, street, name, rating)

Connecting the source schemas to the global schema: Global-as-view (GAV)

s1:ParisHotels(street, name, roomPrice)

s2:LyonHotel(street, name, roomDesc, roomPrice)

s3:Restaurant(city, street, name, rating)

Global: Hotel(city, street, name, descr, price),
Restaurant(city, street, name, rating)

Defining Hotel as a view over the source schemas:

define view Hotel as

select 'Paris' as city, street, name, null as roomDesc, roomPrice as price
from s1:ParisHotels

union all

select 'Lyon' as city, street, name, descr as roomDesc, price
from s2:LyonHotel

Defining Restaurant as a view over the source schemas:

define view Restaurant as select * from s3:Restaurant

Query processing in global-as-view (GAV)

```
define view Hotel as  
select 'Paris' as city, street, name, null as roomDesc, roomPrice as price  
from s1:ParisHotels  
union all  
select 'Lyon' as city, street, name, descr as roomDesc, price  
from s2:LyonHotel
```

Query:

```
select * from Hotel where city='Paris' and price<200 →  
select * from (select 'Paris' as city... union... select 'Lyon' as city...) where  
city='Paris' and price < 200 →  
select * from (select 'Paris' as city...) where city='Paris' and price < 200 →  
select * from s1:ParisHotels where price < 200
```

Query processing in global-as-view (GAV)

```
define view Hotel as
select 'Paris' as city, street, name, null as roomDesc, roomPrice as price
from s1:ParisHotels
union all
select 'Lyon' as city, street, name, descr as roomDesc, price from s2:LyonHotel
define view Restaurant as select * from s3:Restaurant
```

Query:

```
select h.street, r.rating from Hotel h, Restaurant r where h.city=r.city and
r.city='Lyon' and h.street=r.street and h.price<200 →
select h.street, r.rating from (select 'Paris' as city... from s1:ParisHotels
union all select 'Lyon' as city... from s2:LyonHotel) h, (select * from s3:Restaurant) r
where h.city=r.city and r.city='Lyon' and h.street=r.street and h.price<200 →
select h.street,r.rating from (select ... from s2:LyonHotel) h, s3:Restaurant r where
r.city='Lyon' and h.street=r.street and h.price<200 →
select h.street, r.rating from s2:LyonHotel h, s3:Restaurant r where r.city='Lyon' and
h.price<200 and h.street=r.street
```

Concluding remarks on global-as-view (GLAV)

- Query processing = view unfolding: replacing the view name with its definition and working out simple equivalences from there
 - Allows to push to each data source as much as it can do (trusted heuristic)
- Weakness: changes in the data sources require changes of the global schema and, in the worst case, of all applications written based on this global schema
 - E.g. ,if s4:GrenobleHotel joins the system
 - E.g., if s2:LyonHotel withdraws from the system

Connecting the source schemas to the global schema: Local-as-view (GAV)

s1:ParisHotels(street, name, roomPrice)

s2:LyonHotel(street, name, roomDesc, roomPrice)

s3:Restaurant(city, street, name, rating)

Global: Hotel(city, street, name, descr, price), Restaurant(city, street, name, rating)

Defining s1:ParisHotels as a view over the global schema:

```
define view s1:ParisHotels as  
select street, name, price as roomPrice  
from Hotel where city='Paris'
```

Defining s2:LyonHotel as a view over the global schema:

```
define view s2:LyonHotel as  
select street, name, descr as roomDesc, price as roomPrice  
from Hotel where city='Lyon'
```

Defining s3:Restaurant as a view over the global schema:

```
define view s3:Restaurant as  
select * from Restaurant
```

Query processing in Local-as-View (LAV)

```
define view s1:ParisHotels as  
select street, name, price as roomPrice  
from Hotel where city='Paris'  
define view s2:LyonHotel as  
select street, name, descr as roomDesc, price as  
roomPrice from Hotel where city='Lyon'  
define view s3:Restaurant as  
select * from Restaurant
```

Query:

```
select h.street, h.price, r.rating from Hotel h, Restaurant r  
where r.city=h.city and h.street=r.street
```

Query processing in Local-as-View (LAV)

```
define view s1:ParisHotels as  
select street, name, price as roomPrice  
from Hotel where city='Paris'
```

```
define view s2:LyonHotel as  
select street, name, descr as roomDesc, price as  
roomPrice from Hotel where city='Lyon'
```

```
define view s3:Restaurant as  
select * from Restaurant
```

Step 1: identify
potentially useful
views

Query:

```
select h.street, h.price, r.rating from Hotel h, Restaurant r  
where r.city=h.city and h.street=r.street
```

Query processing in Local-as-View (LAV)

- **Query:**
select h.street, h.price, r.rating from Hotel h, Restaurant r where
r.city=h.city and h.street=r.street
- **Step 2:** generate view combinations that may be used to answer the query
 - s1:ParisHotels and s3:Restaurant
 - s2:LyonHotels and s3:Restaurant
- **Step 3:** for each view combination and each view, check:
 - If the view returns enough attributes: we need the attributes returned by the query and those on which possible query joins are based
 - If the view selections (if any) are compatible with those of the query
- **Step 4:** for each view combination, add the necessary joins among the views, possibly selections and projections → rewriting
- **Step 5:** return the union of the rewritings thus obtained

Query processing in Local-as-View (LAV)

```
define view s1:ParisHotels as... from Hotel where city='Paris'  
define view s2:LyonHotel as... from Hotel where city='Lyon'  
define view s3:Restaurant as select * from Restaurant
```

Query:

```
select h.street, h.price, r.rating from Hotel h, Restaurant r where  
r.city=h.city and h.street=r.street
```

Rewriting of the query using the views:

```
select h1.street, h1.price, r3.rating  
from s1:ParisHotels h1, s3:Restaurant r3  
where h1.city=r3.city and h1.street=r3.street  
union all  
select h2.street, h2.price, r3.rating  
from s2:LyonHotels h2, s3:Restaurant r3  
where h2.city=r3.city and h2.street=r3.street
```

Concluding remarks on Local-as-View (LAV)

Query processing

- The problem of finding all rewritings given the source and global schemas and the view definitions is NP-hard in the size of the (schema+view definitions)
- These are still often much smaller than the data
- Fundamental concept: containment mappings between a view and the query [Chandra and Merlin, 1978]

The schema definition is more robust:

- One can independently add/remove sources from the system without the global schema being affected at all
- Thus, no application needs to be aware of the changes in the schema